

ENTR 3360 Assignment 1: Discovery and Product Definition

Project: Mazbien AI Training Needs Assessment and Solution Designer

Problem Statement

Mazbien is an AI training consulting firm. Its consultants meet with client organizations and design role-specific AI training programs. The main user is the Mazbien consultant or trainer. The project owner confirmed this directly. The job to be done is to help the consultant move from an initial client conversation to a structured, role-specific training recommendation more efficiently.

The current process has a gap between the discovery conversation and the final recommendation. The project owner explained that a consultant needs to work through several stages: discovery information, roles and AI readiness, capability gaps and priority needs, a recommended training pathway, package and add-on recommendations, and a draft summary. Only after all of that does the consultant review the work. This chain of reasoning is done manually and inconsistently at this point. The project owner said recommendations need to fit organizations with different employee roles and different levels of AI readiness "rather than providing one generic recommendation for everyone." This suggests the current process does not always differentiate at that level.

This inconsistency costs consultants meaningful time per client, and similar clients sometimes get different recommendations just because a different consultant handled the call. This is based on the project owner's focus on consistency and scalability, not on a specific number or example they gave.

The confirmed goal is to make Mazbien's assessment and recommendation process "faster, more consistent, measurable, and scalable, while still retaining the consultant's professional judgment." The project owner also redirected the design away from any automatic finalization step. This confirms the problem is not just that consultants take too long. It is that they take too long to reach a recommendation they still need to personally stand behind.

Call Evidence

Call prep brief →

Before I met with the project owner, I wanted to answer these open questions:

Is the primary user the consultant or the client organization?

What order should the recommendation logic follow? Specifically, is AI readiness scored before or after capability gaps are identified?

How are package tiers and add-ons actually decided? Are they fixed by department, or based on other factors like organization size?

Should the tool be able to finalize or send a recommendation on its own, or is a human step required?

The project owner sent a written recap before the build continued. These were the confirmed points:

Proposed workflow: Discovery information, Roles and AI readiness, Capability gaps and priority needs, Recommended training pathway, Package and add-on recommendations, Draft recommendation summary, Consultant review.

The primary user is the Mazbien consultant or trainer. The solution should help the consultant move more efficiently from an initial client conversation to a structured, role-specific training recommendation.

The tool must fit organizations with different employee roles and different levels of AI readiness. It should not give one generic recommendation for everyone.

Service tiers are influenced by the specific offerings involved. Add-ons are based on factors like organization size and specific needs, not a fixed department mapping.

Recommendations must be explainable. The consultant needs to understand why a particular training pathway, package, or add-on was suggested.

After reviewing, the project owner also mentioned that the consultant must remain in control of the final recommendation and be able to review or adjust the system's suggestion. The workflow should not end in an immediate payment step. The final action should be a "Send For Review" button.

What changed in my understanding because of this

Unfortunately, I was unable to join the meeting, but after reviewing the email, two things changed for me. First, I had been treating tier and add-on assignment as something mostly determined by which departments were selected. The project owner's note that add-ons are driven by size and specific needs, not a fixed department map, meant the scoring logic needed more inputs than department selection alone. Second, I had originally scoped the workflow to end in an approval-and-send action or even a payment action. The project owner's explicit reversal changed the MVP boundary itself. The system's output is now a draft that the consultant reviews and can adjust. It is never a final client-facing artifact that the tool produces or sends on its own. This is reflected in the MVP scope and workflows below.

A separate Potential Customer Call has not been held so far. The project owner recap above is the main confirmed evidence this submission is based on.

Feature List: MVP vs. Out of Scope

MVP

- Structured client intake that captures organization info, departments, and specific employee roles. This is directly required to accommodate different roles instead of giving just one recommendation.
- Adaptive follow-up questions generated from gaps in the intake. This supports faster and more complete discovery without turning the process into a fully open-ended interview.
- AI-readiness scoring per role or department across defined axes, shown visually. This makes "Roles and AI readiness" its own confirmed workflow stage instead of folding it into one generic score.
- Capability gap and priority need identification per role. This is the confirmed stage between readiness scoring and pathway recommendation.
- Recommended training pathway generated per role. This is required by the explicit role-specific framing, not just department-specific framing.
- Rule-based package and tier mapping, with add-ons driven by organization size and specific needs. This matches the confirmed logic that offerings, not just department, drive tiers and add-ons.
- Visible rationale attached to every generated recommendation, such as "why this tier" and "why this pathway."
- Consultant review and edit screen for both the client-facing recommendation and the internal delivery plan. This is required to keep the consultant in control, based on the project owner's reversal.
- A "Send For Review" action as the final step. This was explicitly specified by the project owner instead of any auto-finalize or payment step.
- Assessment history, so a consultant can reference how similar roles were assessed before.

Out of scope

- Any in-app payment or auto-send-to-client step. The project owner explicitly ruled this out by saying "instead of allowing the user to pay right away." Including it would directly contradict a confirmed requirement.
- A fully open-ended, free-form AI-driven interview instead of guided follow-up questions. This would be harder to keep explainable and consistent across similar clients, which conflicts with the confirmed consistency and explainability goals.
- Auto-generated Premium or enterprise-tier recommendations. That scope is custom by nature. Auto-generating it would create false confidence in a recommendation that actually needs manual scoping and would undermine consultant judgment. The MVP flags Premium for manual configuration instead.
- Multi-consultant collaboration or permissions. This adds access-control complexity that is not required to serve the single confirmed primary user, who is an individual consultant.

- An override or rationale audit log and an "AI accuracy over time" analytics dashboard. These are valuable later, but they depend on real usage data from the review flow existing first. They are premature before the core loop is validated.
- A client-facing self-serve portal. The confirmed primary user is the consultant, not the client. A client-facing product is a different job to be done than the one confirmed.

Critical Workflows

Workflow A: Discovery to draft recommendation

1. Consultant starts a new assessment and enters company name, industry, and size.
2. Consultant enters the departments and specific roles gathered from the discovery conversation.
3. Consultant enters current tools and workflows, pain points, and the client's AI experience level.
4. System reviews the entered data and surfaces a short set of adaptive follow-up questions targeted at the specific gaps in this intake. For example, it might ask about ticket volume if Customer Service is selected, or about change-management concerns if AI experience is "None."
5. Consultant answers what the client already told them during the call, or leaves a question blank to revisit live.
6. Consultant selects "Generate assessment."
7. System scores AI readiness per role or department across the defined axes and displays it visually.
8. System identifies capability gaps and priority needs per role, based on the readiness scores.
9. System generates a recommended training pathway per role, maps the client to a package tier, and suggests add-ons based on size and stated needs. Each item is shown with a visible rationale.
10. Consultant lands on the review dashboard with all of the above populated.

Workflow B: Consultant review and Send For Review

1. Consultant opens the generated dashboard and views both the client-facing recommendation and the internal delivery plan.
2. Consultant reads the rationale behind the assigned tier and each role's training pathway.
3. Consultant edits any field directly where their judgment differs from the system's output. For example, they might adjust a recommended module or manually flag Premium.
4. Consultant optionally adds a short note explaining any change they made.
5. Consultant selects "Send For Review." This finalizes the draft internally but does not trigger billing or send anything to the client automatically.
6. System timestamps and saves the assessment with status "Sent for review." This makes it available in the consultant's history for future similar assessments.

Success Criteria

1. A consultant can go from a completed intake to a reviewable draft recommendation in minutes instead of the current multi-hour manual process. This is measured as time from intake completion to first draft being available.
2. Given two clients with comparable size, department mix, and AI experience, the system assigns the same package tier both times. This is measured by re-running the assessment with near-identical inputs and confirming identical tier output. This directly tests the confirmed consistency goal.
3. For every generated recommendation, a consultant can identify which specific input or inputs produced a given tier, pathway, or add-on without asking anyone else. This is measured by checking that a rationale trace is present and legible for 100% of generated recommendations. This directly tests the confirmed explainability requirement.
4. A consultant can complete a full assessment, from intake through "Send For Review," on their first attempt without external help. This is measured through a short usability test that tracks completion and whether they needed to ask a question.
5. No recommendation reaches "Sent for review" status without passing through the consultant's review screen. This is measured by confirming there is no path in the product that skips this step. This directly tests whether consultant judgment is actually retained, not just implied.

Primary User and MVP Boundary

Primary user: the Mazbien consultant or trainer. The solution is designed around this person's workflow from initial client conversation to structured recommendation. It is not designed around the client organization filling out a form.

MVP boundary: one client organization is assessed at a time, with a defined cap on departments and roles per assessment. For example, up to five departments and ten distinct roles. Package tiers are limited to three deterministic, rule-based tiers: Bronze, Silver, and Gold. Premium is explicitly excluded from auto-generation and is instead flagged for manual configuration, since enterprise-level scope cannot be responsibly reduced to a rule. Every recommendation carries a visible rationale. Nothing is finalized without passing through the consultant's review screen and an explicit "Send For Review" action. There is no in-app payment or auto-send step in this MVP.

What makes this build genuinely mine, given that other students are working from the same project brief: the product is designed around the consultant's review and adjustment experience as the central screen, not the client's intake form. Tier and add-on logic are treated as fully rule-based and are always paired with a visible rationale instead of opaque model output. This makes the explainability requirement actually verifiable. Training pathways are generated at the individual role level, not just the department level. This follows from the project owner's explicit "different employee roles" requirement, but it is more precise than the brief's wording strictly forces. A different student could reasonably have stopped at department-level recommendations instead.

Visuals

Mazbien

ENTR 3360
SOLUTION
DESIGNER

AI Training Needs Assessment & Proposal Engineering

Discovery Intake

Consultant Review & Edit

Draft

Assessment History

Consistency Lab (SC #9)

Quick Demo Presets...

Elena Rostova
(Lead Consultant)

1 Organization Profile

2 Departments & Roles

3 Tools & AI Maturity

4 Adaptive Follow-Ups

Step 1: Client Organization & Contact Information

DISCOVERY PHASE

Capture foundational client profile gathered during the initial discovery call.

Client Company Name *

e.g. Apex Supply Chain Logistics

Official corporate entity name

Primary Client Stakeholder / Title *

e.g. Marcus Vance, VP Operations

Consultant's key contact person and decision maker

Industry & Operating Sector *

Select Industry Sector...

Inform regulatory constraints and baseline tool familiarity

Organization Total Headcount *

Select Headcount Tier...

Drives rule-based package tier mapping and add-on triggers

Consultant Note: You can load a preset scenario from the top bar at any time.

Continue to Departments & Roles -->

Mazbien

ENTR 3360
SOLUTION
DESIGNER

AI Training Needs Assessment & Proposal Engineering

Discovery Intake

Consultant Review & Edit

Draft

Assessment History

Consistency Lab (SC #9)

Quick Demo Presets...

Elena Rostova
(Lead Consultant)

Client-Facing Proposal

Internal Delivery Plan

Readiness & Gaps Deep-Dive

Consultant Adjustments & Overrides

RECOMMENDED SERVICE PACKAGE

GOLD RECOMMENDED

Gold (Enterprise Transformation)

Comprehensive cross-enterprise AI deployment with champions and governance

\$54,000 - \$78,000

(8 - 10 Weeks • 56 - 72 Total Hours)

VISIBLE RATIONALE: WHY THIS TIER WAS RECOMMENDED

✓ Headcount bracket (21 - 100 employees) or department breath (3 departments, 5 roles) requires cross-functional champion enablement.

✓ Multi-cohort structure required to accommodate concurrent departmental streams.

✓ Includes executive governance integration and custom prompt playbooks for all 5 roles.

INCLUDED PACKAGE DELIVERABLES:

• Full Enterprise Role Matrix Curriculum (across all assessed departments)

• Certified "AI Champion" Enablement Track (internal super-users)

• Executive & Legal Governance Workshop on Data Privacy & IP Protection

• Customized IDE / Software Tooling Integrations (Copilot, Cursor, etc.)

DRAFT MODE

All changes saved locally

Re-edit Intake

Send For Review

Mazbien

AI training needs assessment

Client intake

1 Client intake

2 Adaptive follow-up

3 Review & approve

Client intake

Capture the business context. The platform uses this to draft targeted follow-up questions.

Company name

Industry

Company size

e.g. Northfield Logistics

e.g. Freight & distribution

26-100

Departments involved in this assessment

Leadership

Finance

Customer Service

Marketing

Human Resources

Operations

Sales

Roles to assess

Role title, e.g. Support Team Lead

Department

+

Add another role

Current tools & workflows

What software, spreadsheets, or systems does the team use day to day?

Repetitive tasks & pain points

What takes too long, gets duplicated, or causes friction?

Mazbien

ENTR 3360
SOLUTION
DESIGNER

AI Training Needs Assessment & Proposal Engineering

Discovery Intake

Consultant Review & Edit

Draft

Assessment History

Consistency Lab (SC #9)

Quick Demo Presets...

Elena Rostova
(Lead Consultant)

Mazbien Assessment History & Client Benchmarks

Search past clients or industries...

Review previously completed assessments to benchmark recommendations for similar company sizes and department mixes. Demonstrates Mazbien's institutional consistency over time.

Client Organization	Industry	Headcount	Assigned Tier	Status	Date	Action
Apex Supply Chain Logistics	Logistics, Supply Chain & Transport	220 staff	GOLD	SENT FOR REVIEW	Sep 14, 2026	Inspect
FinVantage Wealth Advisory	Finance, Banking & FinTech	45 staff	SILVER	SENT FOR REVIEW	Sep 17, 2026	Inspect
HealthSync Dynamics	Healthcare, Life Sciences & Biotech	850 staff	PREMIUM	SENT FOR REVIEW	Sep 19, 2026	Inspect