

# Mazbien AI Training Needs Assessment — Project Plan

## 1. Primary user & framing

**Primary user:** the Mazbien consultant, not the client organization. The client intake form is the input; the **consultant's review cockpit is the hero screen**. Everyone else in the class will likely design for the client filling out a form — we're designing for the person who has to turn that form into a trustworthy, billable recommendation in minutes instead of hours.

**Job to be done:** given intake + follow-up answers, get to a draft the consultant trusts enough to lightly edit and approve, with visible reasoning behind every recommendation and every price.

## 2. Baseline scope (what every team will build)

- Intake form
- AI-generated follow-up questionnaire
- Recommendation report with tier assignment (Bronze / Silver / Gold / Premium)
- A consultant review screen

We need all of this working end-to-end first. Everything in Section 3 is what we layer on top once the baseline flow is solid.

### 3. Feature list & current status

| #  | Feature                                                                                              | Status                   | Notes                                                                                                   |
|----|------------------------------------------------------------------------------------------------------|--------------------------|---------------------------------------------------------------------------------------------------------|
| 1  | Multi-step client intake (company info, departments, roles/tools, pain points, AI experience, goals) | <b>Built</b>             | 4 sub-steps, validates before advancing                                                                 |
| 2  | Adaptive follow-up questions                                                                         | <b>Built (simulated)</b> | Rule-based generation by department/AI experience/keywords; revealed progressively to simulate analysis |
| 3  | Tier assignment engine (Bronze/Silver/Gold/Premium)                                                  | <b>Built</b>             | Deterministic, driven by org size, department count, AI experience                                      |
| 4  | AI Readiness Radar (5-axis chart per department)                                                     | <b>Built</b>             | Axes: AI Awareness, Tool Fluency, Process Repeatability, Data Literacy, Change Readiness                |
| 5  | Role-based training pathway cards                                                                    | <b>Built</b>             | Per-role, not just per-department; priority-tagged                                                      |
| 6  | Dynamic pricing justifier ("why this tier")                                                          | <b>Built</b>             | Shows which specific answers triggered the tier                                                         |
| 7  | Internal delivery plan tab                                                                           | <b>Built</b>             | Module sequencing, proposal summary, implementation checklist                                           |
| 8  | Consultant notes / override field                                                                    | <b>Built (basic)</b>     | Free-text field exists                                                                                  |
| 9  | Override & rationale <b>log</b> (tracked history, not just a field)                                  | <b>Not yet built</b>     | Needs a persisted record of what changed + why, separate from the live draft                            |
| 10 | "AI accuracy vs. consultant correction" dashboard                                                    | <b>Not yet built</b>     | Depends on #9 existing first                                                                            |
| 11 | Same-day guarantee timer / turnaround tracker                                                        | <b>Not yet built</b>     | Live countdown or status (green/yellow/red) from intake → draft → approval                              |

| #  | Feature                                      | Status               | Notes                                                                    |
|----|----------------------------------------------|----------------------|--------------------------------------------------------------------------|
| 12 | Drag-reorder on role cards                   | <b>Not yet built</b> | Nice-to-have polish once cards are stable                                |
| 13 | Assessment history / reuse for similar roles | <b>Not yet built</b> | Needed to hit the "consistency across similar roles" success target      |
| 14 | Real LLM calls (vs. rule-based simulation)   | <b>Not yet built</b> | Current questions/text are rule-generated for demo speed and determinism |

## 4. Detailed specs for the differentiator features

### 4.1 AI Readiness Radar

- One radar per department (or org-wide overview + drill-down per department).
- 5 axes, each scored 0–5: AI Awareness, Tool Fluency, Process Repeatability, Data Literacy, Change Readiness.
- Scores derived from intake answers + follow-up answers (e.g., "0 AI experience" → low Awareness; "manual spreadsheet reporting" → low Process Repeatability, higher automation opportunity).
- **Open question for team:** do we score at intake time only, or re-score after follow-up answers come in? Recommend re-scoring after follow-up, so the radar reflects the fuller picture.

### 4.2 Dynamic pricing justifier

- Every tier recommendation ships with a visible rule trace: "Your HR team has 0 AI experience + 3 repetitive workflows → Silver minimum."
- Implementation approach: keep tier logic as **explicit rules**, not LLM judgment, and log which rule(s) fired. This is what makes pricing consistent (a stated Mazbien goal) and auditable.
- **Open question:** do we want the justifier to show *all* triggered rules, or just the single highest-weight one? All of them is more transparent but busier — worth a quick usability check.

### 4.3 Role-based training path cards

- One card per named role (not department): title, priority tag, recommended modules, timeline, expected outcome.
- Consultant can reorder before approving (drag-and-drop) — signals this is a working draft, not a locked report.

- **Open question:** does reordering actually change anything downstream (e.g., sequencing in the internal plan), or is it purely presentational for now? Decide before building drag logic — if it's cosmetic only, a simpler up/down control may be enough for MVP.

#### 4.4 Override & rationale log

- Every time a consultant edits AI-generated text, capture: what changed (before/after), and a short free-text reason.
- Store this as a running log per assessment, separate from the current draft state.
- Longer-term payoff: a simple "AI accuracy rate" view (% of sections approved without edits) — a genuinely compelling story for demo day, but only build the dashboard after the log itself is reliable.
- **Open question:** is capturing a reason *required* on every edit, or optional? Required gives better data but adds friction — optional is safer for MVP, can nudge with a placeholder like "why did you change this? (optional)".

#### 4.5 Same-day guarantee timer

- Track timestamps: intake submitted → follow-up generated → draft generated → consultant approved.
- Status color: green (on track for same-day), yellow (approaching end of day), red (missed same-day target).
- This directly visualizes one of Mazbien's stated success metrics, not just a UI flourish.
- **Open question:** what's the actual target window — same calendar day, or a rolling X-hour window? Needs a real number from the brief or from whoever owns the metric, since "same-day" is ambiguous near end of business hours.

### 5. What still needs deciding before the team splits up work

1. Real LLM integration: which calls become live model calls (follow-up question generation, draft recommendation text) vs. staying rule-based for reliability (tier math, pricing logic should stay rule-based regardless).
2. Whether assessment history/reuse (item 13) is in scope for the grading deadline or explicitly called out as a stretch goal in the write-up.
3. Data model for the override log — needs a shape agreed on before anyone builds the accuracy dashboard on top of it.
4. Same-day timer's actual target window (see 4.5).
5. Whether role-card drag-reorder has real downstream effect or is presentational for now (see 4.3).

## 6. Suggested build order

1. Lock the baseline flow end-to-end (intake → follow-up → tier + recommendation → approve) — **done**.
2. Radar chart + pricing justifier + role cards — **done**, review for polish and edge cases (e.g., a department with no mapped use cases).
3. Override & rationale log (item 9) — foundational for the accuracy dashboard, so prioritize this next.
4. Same-day timer (item 11) — self-contained, can be built in parallel by someone else.
5. Assessment history/reuse (item 13) — needed for the consistency success metric.
6. AI accuracy dashboard (item 10) — only after the log has real data flowing through it.
7. Real LLM calls to replace rule-based simulation, where it adds clear value over deterministic rules.
8. Drag-reorder polish on role cards, if decided to be more than cosmetic.

## 7. Tie-back to Mazbien's stated success targets

| Mazbien target                             | How this build addresses it                                                                                |
|--------------------------------------------|------------------------------------------------------------------------------------------------------------|
| Cut consultant analysis time<br>~40%       | Draft recommendation + radar + role cards generated automatically from intake/follow-up, not built by hand |
| Same-business-day drafts                   | Same-day timer makes turnaround visible and trackable, not just assumed                                    |
| Consistent pricing for<br>comparable scope | Tier logic is rule-based and its rationale is shown, not left to individual judgment                       |
| Systematic add-on/upsell<br>surfacing      | Add-on suggestions generated from defined triggers tied to intake/pain points                              |
| System improves over time                  | Override & rationale log + accuracy dashboard gives a concrete "the AI gets better" story                  |